

sheet05

November 19, 2017

1 Weighted K-Means

In this exercise we will simulate finding good locations for production plants of a company in order to minimize its logistical costs. In particular, we would like to place production plants near customers so as to reduce shipping costs and delivery time.

We assume that the probability of someone being a customer is independent of its geographical location and that the overall cost of delivering products to customers is proportional to the squared Euclidean distance to the closest production plant. Under these assumptions, the K-Means algorithm is an appropriate method to find a good set of locations. Indeed, K-Means finds a spatial clustering of potential customers and the centroid of each cluster can be chosen to be the location of the plant.

Because there are potentially millions of customers, and that it is not scalable to model each customer as a data point in the K-Means procedure, we consider instead as many points as there are geographical locations, and assign to each geographical location a weight w_i corresponding to the number of inhabitants at that location. The resulting problem becomes a weighted version of K-Means where we seek to minimize the objective:

$$J(c_1, \dots, c_K) = \frac{\sum_i w_i \min_k \|x_i - c_k\|^2}{\sum_i w_i},$$

where c_k is the k th centroid, and w_i is the weight of each geographical coordinate x_i . In order to minimize this cost function, we iteratively perform the following EM computations:

- **Expectation step:** Compute the set of points associated to each centroid:

$$\forall 1 \leq k \leq K : \quad \mathcal{C}(k) \leftarrow \left\{ i : k = \arg \min_k \|x_i - c_k\|^2 \right\}$$

- **Minimization step:** Recompute the centroid as a the (weighted) mean of the associated data points:

$$\forall 1 \leq k \leq K : \quad c_k \leftarrow \frac{\sum_{i \in \mathcal{C}(k)} w_i \cdot x_i}{\sum_{i \in \mathcal{C}(k)} w_i}$$

until the objective $J(c_1, \dots, c_K)$ has converged.

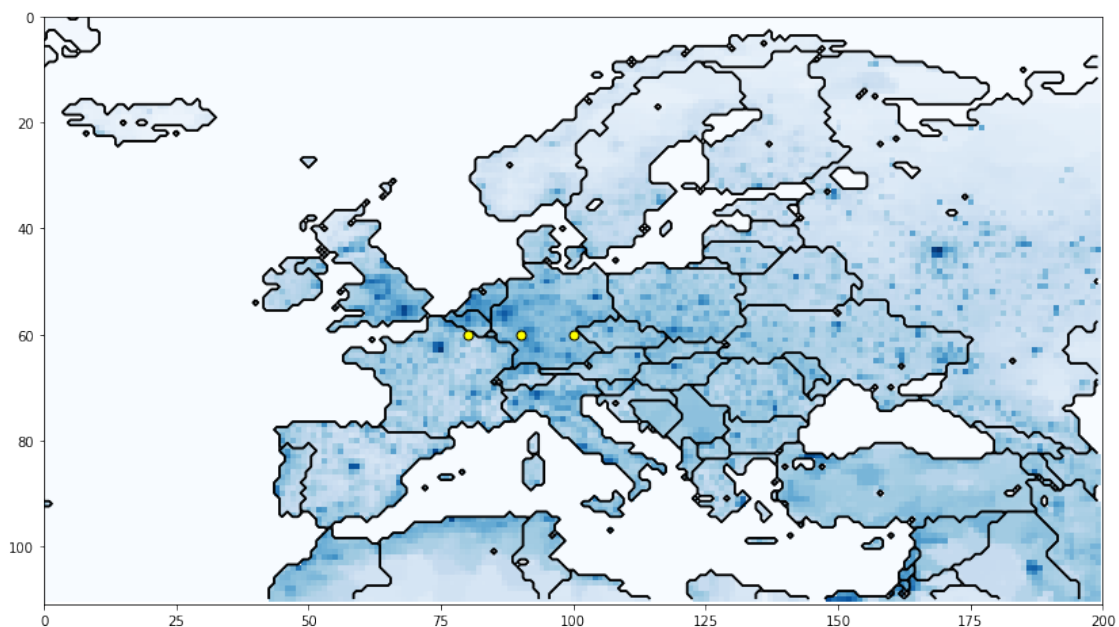
1.1 Getting started

In this exercise we will use data from <http://sedac.ciesin.columbia.edu/>, that we store in the files `data.mat` as part of the zip archive. The data contains for each geographical coordinates (latitude and longitude), the number of inhabitants and the corresponding country. Several variables and methods are provided in the file `utils.py`:

- `utils.population` A 2D array with the number of inhabitants at each latitude/longitude.
- `utils.countries` A 2D array with the country indicator at each latitude/longitude.
- `utils.nx` The number of latitudes considered.
- `utils.ny` The number of longitudes considered.
- `utils.plot(latitudes,longitudes)` Plot a list of centroids given as geographical coordinates in overlay to the population density map.

The code below plots three factories (white squares) with geographical coordinates (60,80), (60,90),(60,100) given as input.

```
In [1]: import utils
        %matplotlib inline
        utils.plot([60,60,60],[80,90,100])
```



1.2 Initializing Weighted K-Means (15 P)

Because K-means has a non-convex objective, choosing a good initial set of centroids is important. Centroids are drawn from the following discrete probability distribution:

$$P(x,y) = \frac{1}{Z} \cdot \text{population}(x,y)$$

where Z is a normalization constant. Furthermore, to avoid identical centroids, we add a small Gaussian noise to the location of centroids, with standard deviation 0.01.

Tasks:

- **Implement the initialization procedure above.**
- **Run the initialization procedure for $K=200$ clusters.**
- **Visualize the centroids obtained with your initialization procedure using `utils.plot`.**

```
In [2]: import numpy as np
```

```
population = utils.population
```

```
# Initializing weighted K means
```

```
def initial(population, sd = 0.01, K = 200):
```

```
    P = population/np.sum(population)
```

```
    initial_centroids = np.random.choice(range(0,P.size), size = K, replace = False, p
```

```
    location = np.unravel_index(initial_centroids, P.shape)
```

```
    longitude = list(location[0]) + np.random.normal(scale = sd, size = location[0].sha
```

```
    latitude = list(location[1]) + np.random.normal(scale = sd, size = location[1].sha
```

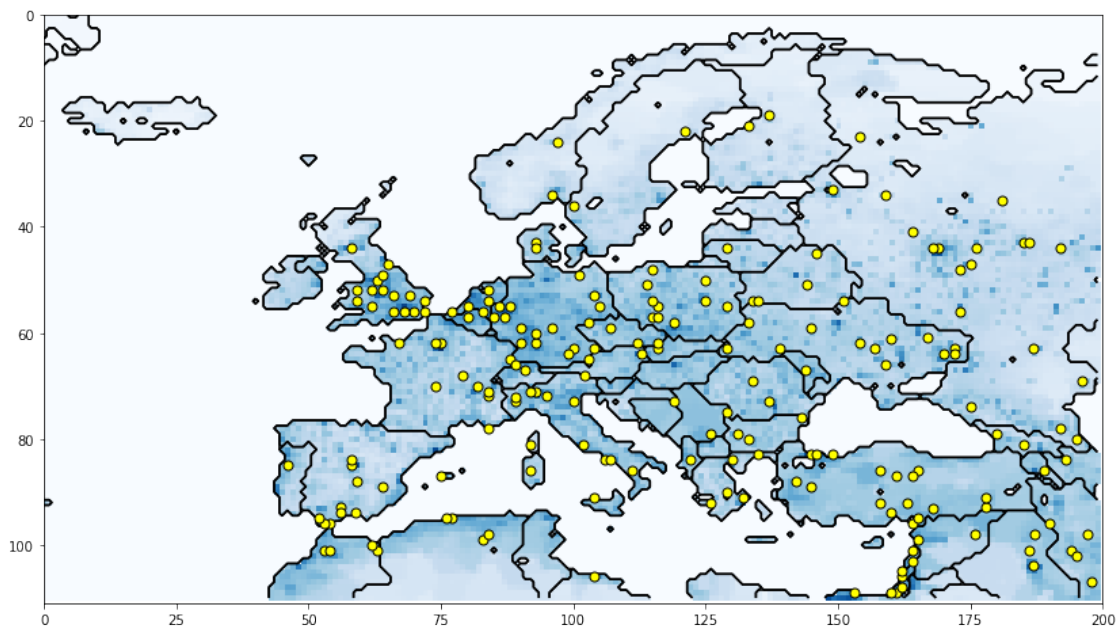
```
    return np.transpose([longitude, latitude])
```

```
# Generate initial centroids
```

```
centroids_1 = initial(population)
```

```
# Plot initial centroids
```

```
utils.plot(centroids_1[:,0], centroids_1[:,1])
```



1.3 Implementing Weighted K-Means (30 P)

Tasks:

- Implement the weighted K-Means algorithm as described in the introduction.
- Run the algorithm with $K=200$ centroids until convergence (stop if the objective does not improve by more than 0.01). Convergence should occur after less than 50 iterations. If it takes longer, something must be wrong.
- Print the value of the objective function at each iteration.
- Visualize the centroids at the end of the training procedure using the methods `utils.plot`.

```
In [3]: def J(centroids,population):
        w = (float(1)/population.sum())*population
        Lon,Lat = population.shape
        J = 0
        for i in range(0,Lon):
            for j in range(0,Lat):
                if w[i,j] > 0:
                    J = J + w[i,j]*min((i - centroids[:,0])**2 + (j - centroids[:,1])**2)
        return J

def kmeans(population, improve, K = 200):
    # Initial centroids
    centroids = initial(population)

    # Initial J functon
    J_cost = J(centroids, population)

    # Some parameters
    delta = 1
    iteration = 0

    Lon,Lat = population.shape

    while delta > improve:

        ## Expectation step
        cluster_assignment = np.zeros((Lon,Lat))
        for i in range(0,Lon):
            for j in range(0,Lat):
                if population[i,j] > 0:
                    cluster_assignment[i,j] = np.argmin((i - centroids[:,0])**2 + (j -
```

```

## Minimization step
centroids = np.zeros((K,2))
for k in range(K):
    k_cluster = np.multiply(np.equal(cluster_assignment,k)*1, population)
    w = (float(1)/k_cluster.sum())*k_cluster
    for i in range(0,Lon):
        for j in range(0,Lat):
            if k_cluster[i,j] > 0:
                centroids[k,:] = centroids[k,:] + w[i,j] * np.array([i,j])

# New J cost function
New_J = J(centroids, population)

# Difference between actual J and previous J
delta = J_cost - New_J

# Overwrite J value
J_cost = New_J

# Count iteration and print result
iteration = iteration + 1
print("{:2d}: J = {:.4f}".format(iteration, J_cost))

return centroids

# kmeans algorithm with plot
centroids = kmeans(utils.population, improve = 0.01)
utils.plot(centroids[:,0],centroids[:,1])

```

```

1: J = 11.15
2: J = 9.82
3: J = 9.26
4: J = 8.78
5: J = 8.40
6: J = 8.14
7: J = 7.95
8: J = 7.74
9: J = 7.66
10: J = 7.61
11: J = 7.57
12: J = 7.52
13: J = 7.49
14: J = 7.47
15: J = 7.45
16: J = 7.44
17: J = 7.43
18: J = 7.42

```

19: $J = 7.41$

